

Final Project: An LC-3 Simulator

CS 350: Computer Organization & Assembler Language Programming

Phase 1: Due Sat Nov 19; Phase 2: Due Sat Dec 3

A. Why?

- Implementing a computer really helps you understand its structure.

B. Outcomes

After this project, you should

- Know how to implement the major parts of a computer in a higher-level language.
- Know how to access structures via pointers in C.

C. Project

- You are to implement of a simple console-oriented simulator for the LC-3 in C using your SDC simulator from the labs as a starting point.
- Your program should take the name of a *.hex file as a command line parameter, load it into memory, initialize the control unit, and run a command loop. Commands can cause program execution and other functions.
- Phase 1 of the project only needs to load the hex file into memory; Phase 2 includes handling the simulator commands and executing instructions.

D. Simulating The LC-3 Architecture

- As in the SDC simulator, the LC-3 CPU should be modeled as a value of type `struct CPU`, allocated in the main program. Routines that need the CPU should be passed a pointer to the CPU value.
- A `Word` is `typedef`'ed to be a `short int`, 16 bits on the `fourier` machine. An `Address` is also 16 bits but it is an `unsigned short int` (compare with the `unsigned char` used for the SDC).

- Represent the CPU's memory as an array of `Word` values indexed by `Address` values. Represent the general purpose registers as an array of `Word` values indexed by register numbers (0 – 7). The PC is an `Address`; the IR is a `Word`.
- Below is a skeleton declaration for the CPU structure. Rather than mandate that you represent (e.g.) the condition code with a certain datatype, I've used `typedef` to give a name for the datatypes `ConditionCode`, `Opcode`, `Flag`, and `Reg`. You can just replace the `???` with whatever you decide to use (like `char` or `int` etc).

```
// CPU Declarations -- a CPU is a structure with fields for the
// different parts of the CPU.
//
typedef short int Word;           // A word of LC-3 memory
typedef unsigned short int Address; // An LC-3 address
typedef ??? ConditionCode;       // Condition code: <, =, or > 0
typedef ??? Opcode;              // Opcode 0-15
typedef ??? Flag;                // Boolean flag
typedef ??? Reg;                 // Register number 0-7

#define MEMLEN 65536
#define NREG 8

typedef struct {
    Word mem[MEMLEN];
    Word reg[NREG];              // Note: "register" is a reserved word
    Address pc;                  // Program Counter
    ConditionCode cc;            // positive, zero, or negative
    Flag running;                // is CPU executing instructions?
    Word ir;                     // Instruction Register
    Opcode opcode;               // opcode 0-15
} CPU;
```

- In real life, you'd put the type definitions and similar items in a header file (`*.h`), but for this project, don't bother.

E. Interactions Between Signed/Unsigned and Short/regular integers

- Here are some warnings about the interactions between signed and unsigned and short or regular integers:
- Since unsigned values are never negative, as `Address` values, `0x8000 – 0xffff` represent 32768 to 65535. Since `Word` values are signed, `0x8000 - 0xffff` represent -32768 to -1.

- The arithmetic operators (+, −, etc.) yield regular-sized integer results, not `short` integers. If you want to cast a regular integer into a `Word` or `Address`, you can use a cast expression `(Word) regular_integer` or `(Address) regular_integer`.
- Similarly, if you need to take a `Word` value and treat it as an `Address`, you can cast it using `(Address) word_value`. (This is helpful for converting base register values to addresses in `LDR` and `STR`.)
- Loops of the form “For every `Address`, do something” are a little tricky.
 - If `i` is an `Address` variable, then `for (i = 0; i < 65536; i++) { ... }` will cause an infinite loop because an `Address` is always less than 65536. (The `gcc` compiler will give you a warning — heed it.)
 - Various alternatives:
 - You can check for `i == 65535` (but you'll be one iteration short).
 - You can check for the **second** time that `i == 0` (which is a little tricky).
 - You can make `i` not be `short`.

F. Phase 1: Everything Except for Actual Instruction Execution

- As in the SDC simulator, the input file should be a command line parameter. If it's omitted, use `default.hex` as the default filename.
- You should be able to process a `*.hex` file produced by the LC-3 editor when it assembles a `*.asm` file.
 - A `*.hex` file is a text file containing a sequence of lines with a four-digit hex number on each line (no leading `x` or `0x`). An `scanf` format of `%x` will read in an unsigned integer, which you can cast to a `Word` of memory.
 - The first line specifies the `.ORIG` location to start loading values into. The remaining lines contain the values to load.
 - If you read a value into location `xFFFF`, wrap around to `x0000` as the next location. No complaining is necessary.
 - If anything appears on a line after the four-digit hex number, ignore it. (This will let us add comments to our `hex` files.)

- **Note:** In general, the LC-3 text editor can translate hex files to executable object code via *Translate* → *Convert Base 16*, but it won't handle a hex file with comments added to it.
- All other memory locations should be set to the value zero. Don't simulate the TRAP table in low memory (x00 – xFF) or the trap-handling code or memory-mapped I/O addresses in high memory.
- Once the program is read into memory, initialize the PC to the .ORIG value, the IR to zero, the condition code to Z, and the running flag to true, then dump the CPU and memory.
 - For the CPU, print out the PC, IR, CC, running flag, and data registers.
 - For memory, print out the memory address and its value, but **only for nonzero values**. Don't print 2^{16} lines of memory values (almost all of which are zero).
 - Print out the memory value three times: Once in hex, once in decimal, and once using assembler mnemonics. For example, if locations x8000 – x8002 contain x12A0, 0, and x172F respectively, you might print


```
x8000: x12A0    4768  ADD    R1, R2, 0
x8002: x172F    5935  ADD    R3, R4, 15
```
 - Note that location x8001 wasn't printed out, since it contains a zero.
- Once you do the initial dump, pass control to the simulator command loop. As with the SDC simulator, you should repeatedly prompt for a command, read it in, and execute it. For Phase 1, don't actually execute the LC-3 instructions; you'll add that in Phase 2.

Command	Syntax	Action
Help	? or h	Print a summary of the simulator commands
Dump	d	Print the CPU and memory addresses/values
Execute	<i>integer</i>	Execute that many instruction cycles
Execute	(none: just a \n)	Equivalent to the integer 1 (execute one cycle)
Quit	q	End the simulation

Simulator Commands

- The commands and command format are the same as in the SDC simulator: They should appear one per line, where each line is terminated by a '\n'.
- You should ignore white space (spaces or tabs) before or after the command. (E.g., “ 12 \n” and “12\n” should behave the same; so should “\n” and “ \n”.) As with the SDC simulator, reading a line using `fgets` and `sscanf` formats will take care of these requirements.
- Ignore any text (whitespace or non-whitespace) that appears after the command but before the end of the line. Again, `sscanf` should handle this for you.
- You can assume the command names `?`, `h`, `d`, `q` are in lower case.
- For the `d` (dump) command, print out the CPU and **non-zero** memory values.
- As in the SDC simulator, for the numeric execute command, the number of cycles to execute should be a decimal integer ≥ 1 and \leq some sanity check limit like 100. For an insanely large number, complain and substitute a sane number.
- For Phase 1, just print a message instead of executing whatever number of LC-3 instructions.

G. Phase 2: Executing LC-3 Instructions

- For Phase 2, implement execution of the different LC-3 instructions.
- If you're running a number of instruction cycles and execution halts before you run all cycles, stop early and go on to the next command. (The `HALT` trap stops execution but does not stop the command loop.)
- Print trace information as you execute instructions. You should print out as much trace information as the sample solution, but you don't have to follow its exact format.
- For each instruction you execute, print its memory location in hex and print the instruction in hex and using mnemonics.
- E.g., say we're executing the instruction at location `x2468`, with `M[x2468] = x6281`, then we'd print the location `x2468`, its value `x6281`, and its mnemonic representation `LDR R1, R2, 1`

- If an instruction uses or modifies a register or memory location, give the register name or memory location and its value. (Make sure the trace says whether you're using or modifying the value.)
 - For these purposes, the PC is a register when it's used in PC-offset addressing or when you do a jump or branch.
- If an instruction uses a constant (an offset or immediate value, for example), name it.
- If an instruction calculates an intermediate value (such as *address + offset*), name both parts and the result.
- E.g., continuing with the example of `LDR R1, R2, 1`, if `R2 = x1234`, we would show the address calculation `x1235 = x1234 + 1` and the value of `M[x1235]` that we're copying to R1.

H. Differences From the Patt & Patel Simulator

Your simulator should produce the same results as the Patt & Patel simulator with some slight differences.

- If you execute an instruction at `xFFFF` then incrementing the PC should wrap it around to `x0000`. (Patt & Patel's simulator causes an error if you execute the instruction at `xFFFF`.) By default, for us, `M[xFFFF] = 0` is a NOP.
- Executing the RTI instruction (opcode 8: Return From Interrupt) or the unused opcode 13 should print an error message but continue execution. (Patt & Patel's LC-3 simulator behaves differently.)
- Only traps `x20`, `x21`, `x22`, `x23`, and `x25` need to be implemented. For any other trap vector (including `x24`: PUTSP), print an error message and halt execution (set the running flag to false).
- Execute each TRAP command in one instruction cycle. (Don't simulate the I/O registers.) E.g., executing PUTS should print out the entire string pointed to by R0. The different TRAP instructions do set the condition code (see the technical footnote in the **Branch Instruction** part of Lecture 10).

- For the IN and GETC traps, the user should enter a `\n` after the character to be read in. If the user just enters `\n` without a preceding character, then use `\n` as the character read in.
- Technically, the OUT trap is only supposed to print the right byte of R0, but if your OUT only works if the left byte of R0 is `x00` or `xFF`, that's okay. Similarly, the IN and GETC traps are supposed to overwrite only the right byte of R0 and leave the left byte unchanged, but if you sign-fill the character you read in to get 16 bits, that's okay.
- Because the simulator prints out a trace of execution, printing a prompt and doing a read (using PUTS and GETC) doesn't behave exactly like it does with Patt & Patel's simulator: You have to wait until the GETC executes and asks for your input before actually typing in the character.

I. Code Framework

- You should be able to adapt your SDC simulator to do the final project.
- Feel free to use standard library functions like `strcmp`. (Don't forget to `#include <string.h>`.)
- Remember, your program gets tested on `fourier.cs.iit.edu`, so **do not** use nonstandard libraries like `conio`.
- The method for representing the condition code is up to you: an integer is fine, but it could also be the rightmost three bits of an `unsigned char`, for example. Or literally the character `'N'`, `'Z'`, or `'P'`.
- Also feel free to add more fields or named constants to your program, such as `#define CC_ZERO ...` to represent condition code zero.
- When you print a negative `Word` value in hex, just print four hex digits. (So for `-1`, print `xffff`, not `xffffffff`. Hint: Look up the difference between `%x` and `%hx` in `printf` formats.)

J. Sample Solutions and Test Data

- Sample data and solutions are available on `fourier` in the `~sasaki/CS350` directory: The simulators are named `fp1` and `fp2` (for Phase 1 and Phase 2); the

test data is in `fp_testdata.zip`. The `load*.hex` files are intended to test Phase 1; the others are for Phase 2. You should create your own test data to help you test your program as you write it.

K. Due Dates, Collaboration, What to Submit

- Due Dates: Phase 1 is due **Sat Nov 19**; Phase 2 is due **Sat Dec 3**. Phase 1 can be handed in late; Phase 2 can't. For both phases, there's no extension for attending lab.
- As usual, if you submit multiple times, I'll grade the last submission and toss the earlier ones.
- **Collaboration:** Allowed on Phase 1 but not Phase 2. (Working with the TAs or instructor is always okay, however.)
- **What to turn in:** Just submit the `*.c` file to Blackboard. **Don't** zip it, and **don't** include object files or hex files, etc. If your program has bugs or includes extra credit, please add that information to the comments at the top of your `*.c` file.

L. Grading Scheme

Projects will be given a letter grade (with + and - variants possible). Here are descriptions for typical programs for each letter grade.

- A+ Everything for an A plus extra credit.
- A Program is bug-free or has small number of very minor bugs (e.g., spelling errors). Program structure and commenting are well-done: Each routine is cohesive (does just one thing); related routines are collected together and commented. Comments for any tricky code.
- B Program loading and command loop essentially works. Instruction execution has mostly minor bugs, maybe one major bug. Trace output missing a few details or is kind of hard to read. Program structure and commenting okay.
- C Some major bugs (e.g., negative offsets or addresses $\geq \text{x8000}$ fail); some output bugs (e.g., `xffffffff` for `-1`); execution trace missing major pieces (such as readability). Program has some structure issues (e.g., very long routines). Commenting so-so.

- D Compiles without errors or serious warnings. Memory loading works. Command loop mostly okay (e.g., doesn't check for sane limits). Missing/incorrect execution for some instructions. Trace output poor. Program not very structured; no real comments.
- F Compilation causes error or generates serious warning messages.

M.Extra Credit Possibilities

- Extra credit is only available if your program works.
- **Fancy Output.** Some examples:
 - Initial memory dump starts at *origin* and wraps around `xFFFF` to `x0000` to *origin*-1. (The sample solution does this.)
 - Small values are printed in decimal; large ones in hex. (E.g., instead of `x0`, `x1`, ..., some limit, print just 0, 1, 2,
 - Instead of printing "*address* + 0 = *address*", just print "*address*".
- **Extra Commands**
 - Dump a specified section of memory.
 - Change the value of a memory location.
 - Change the value of a register.
 - Change the PC.
 - Change the running flag back to 1 after `HALT` sets it to 0.
 - Change the condition code.